

# **Deep Reinforcement Learning based Recommendation with Explicit User-Item Interactions Modeling**

**Reinforcement Learning final project**

Annalisa Paladino, Luca Pernice, Emanuele Ruoppolo

# Introduction

**Aim**: create a dynamic recommendation system

**Symbiotic** relationship between **users** and the **recommendation** model



The model provides recommendations to users, helping them come to a **decision**, and users provide feedback to the model, helping the model make **better recommendations**

# Non-RL Recommendation Systems

- ❑ Content-based collaborative filtering
- ❑ Matrix factorization based methods
- ❑ Logistic regression
- ❑ Factorization machines and its variants
- ❑ Deep learning models


## Limitations:

1. Consider the recommendation procedure as a **static** process
2. Trained by maximizing the *immediate* rewards of recommendations, but ignores the **long-term** contributions that the items can make

# RL Recommendation Systems

- ❑ **Model-based** (eg. POMDP and Q-learning): inapplicable when the number of candidate items is large, because a time-consuming dynamic programming step is required to update the model

- ❑ **Model-free**



**Value-based:** very inefficient if the action space is too large

**Policy-based:** the interactions between users and items is not explicitly modeled

# Translation into a RL framework

**Agent:** recommendation system

**Environment:** users

**State:** representation of a user and his most recent positively rated items

**Actions:** continuous vector weighting items in a candidate set

**Reward:** user's rating of an item

**Discount factor:** measure the present value of long-term rewards

# Recommender-User interactions in MDP

At timestep  $t$ , the **state** can be defined as

$$s_t = f(H_t)$$

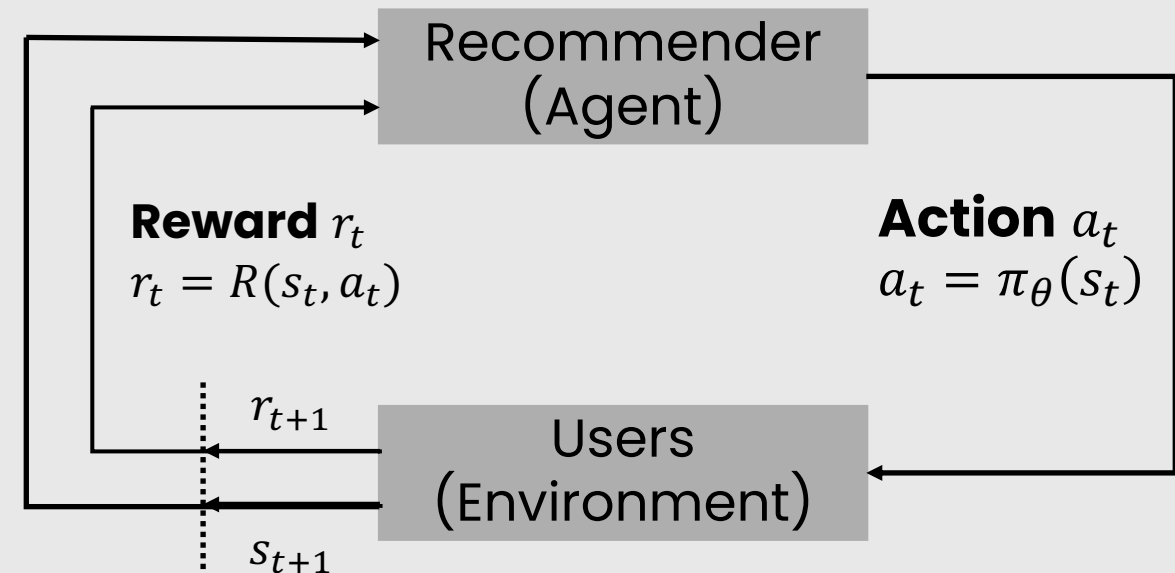
state representation module

$H_t = \{i_1, \dots, i_n\}$  are the embeddings of the latest positive interaction history  
 $i_t$  is a  $k$ -dimensional vector

## Why?

1. A superior recommender system should cater to the users' taste, what items the users like
2. The latest records represent the users' recent interests more precisely.

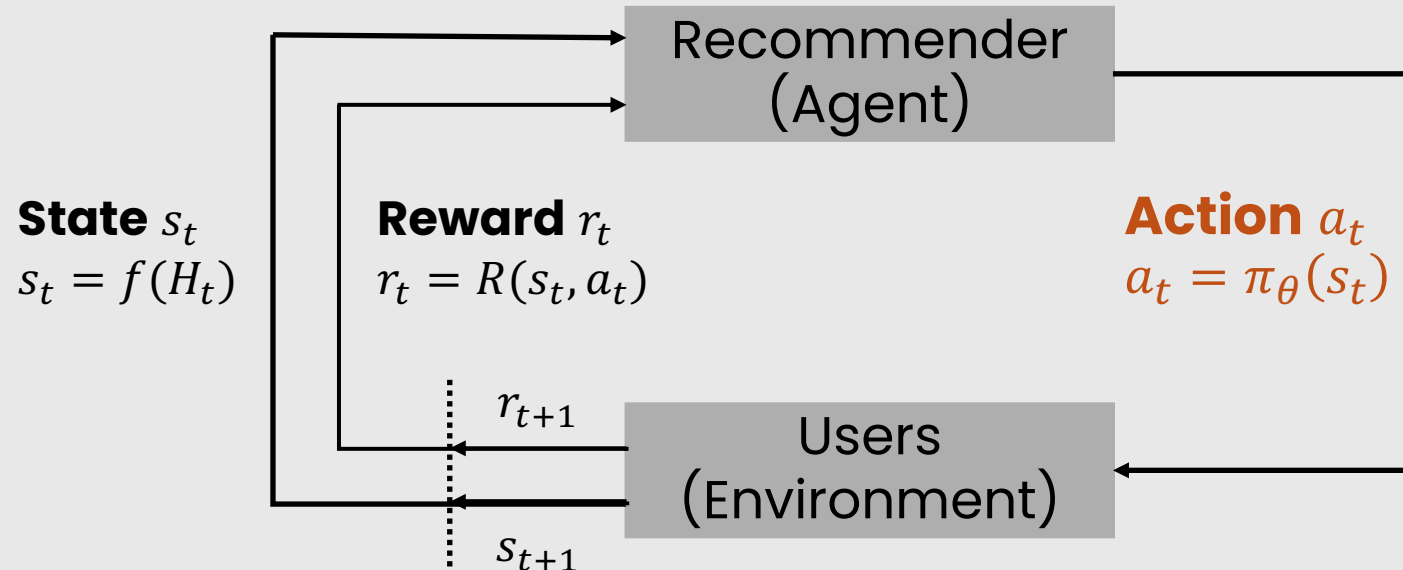
**State**  $s_t$   
 $s_t = f(H_t)$



# Recommender – User interactions in MDP

An **action** is a continuous parameter vector that is used to determine the **ranking scores** of all the items

All the candidate items are ranked according to the computed scores and Top- $N$  items are recommended to the user

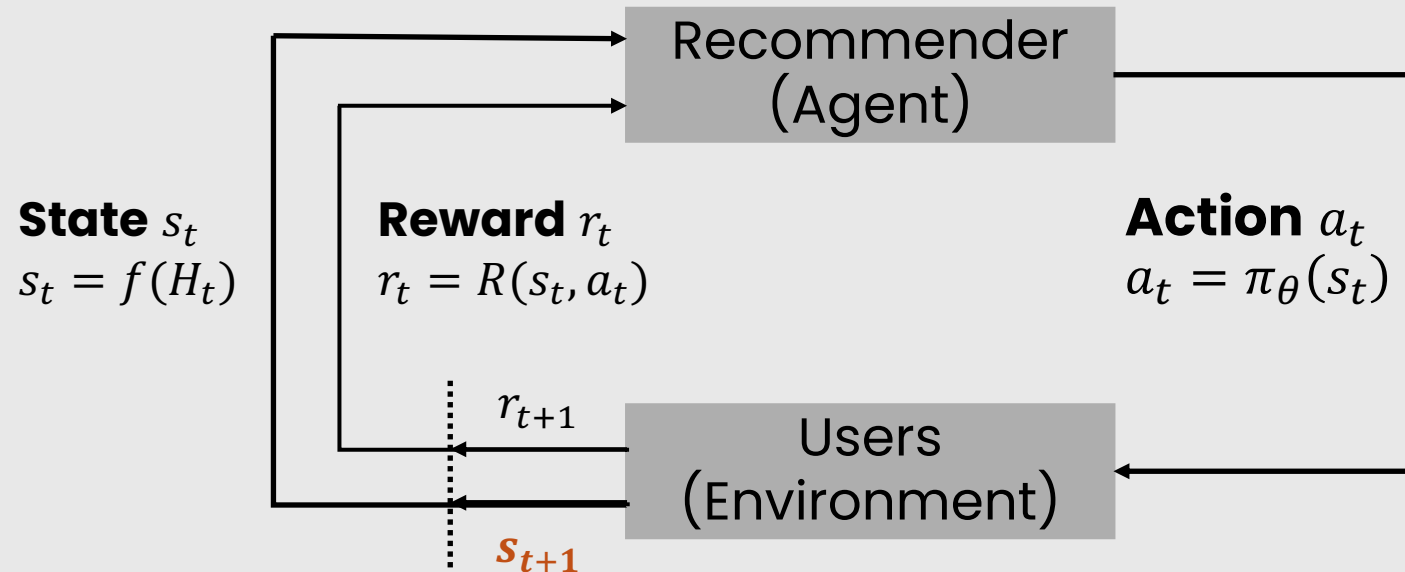


# Recommender–User interactions in MDP

When the recommender agent recommends an item  $i_t$ , if the user provides *positive feedback*, then in the **next timestep**, the state is updated to

$$s_{t+1} = f(H_{t+1})$$

$\longrightarrow H_{t+1} = \{i_2, \dots, i_n, i_t\}$






**BACKGROUND**

# Actor – Critic

## Actor

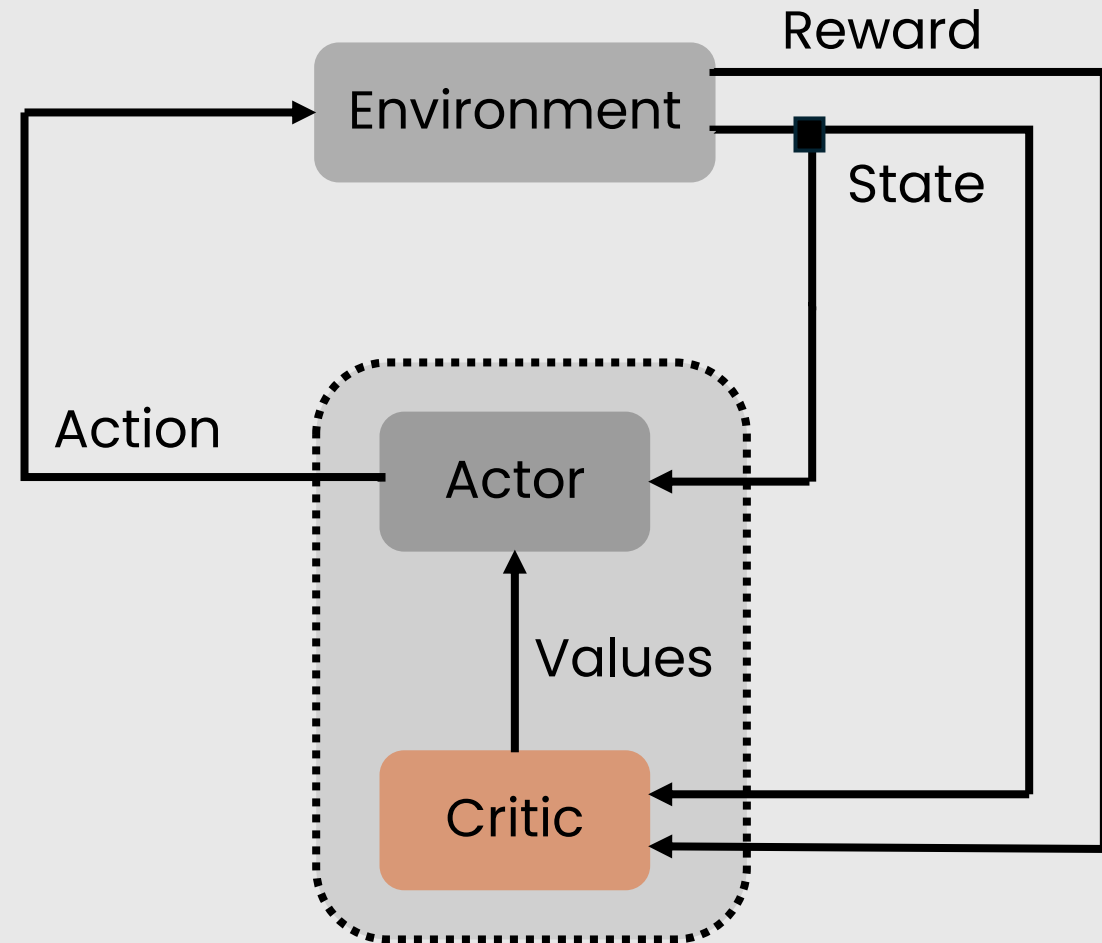
An *agent* (the “actor”) makes decisions by selecting actions based on the current **policy**. It explores the action space to maximize expected cumulative rewards. By continuously refining the policy, the actor adapts to the dynamic nature of the environment



# Actor – Critic

## Critic

A *value function* (the “critic”) evaluates the actions taken by the actor. It estimates the **value** or quality of these actions by providing feedback on their performance. The critic’s role is to guide the actor towards actions that lead to higher expected returns, contributing to the overall improvement of the learning process



# Deep Deterministic Policy Gradient (DDPG)

DDPG is an **off-policy actor-critic** based algorithm

## Key components:

- **Critic Network**  $Q(s, a | \theta^Q)$ : it estimates the action-value function, which predicts the expected return of a given state-action pair
- **Actor Network**  $\mu(s | \theta^\mu)$ : it outputs the action to be taken given a state
- **Target Networks**  $Q'$  and  $\mu'$ : these are copies of the critic and actor networks used for stable updates
- **Replay Buffer**  $R$ : buffer to store experience tuples  $(s_t, a_t, r_t, s_{t+1})$

# DDPG Algorithm

---

**Algorithm** .. DDPG algorithm

---

Randomly initialize critic network  $Q(s, a|\theta^Q)$  and actor  $\mu(s|\theta^\mu)$  with weights  $\theta^Q$  and  $\theta^\mu$ .

Initialize target network  $Q'$  and  $\mu'$  with weights  $\theta^{Q'} \leftarrow \theta^Q$ ,  $\theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer  $R$

**for** episode = 1, M **do**

    Initialize a random process  $\mathcal{N}$  for action exploration

    Receive initial observation state  $s_1$

**for** t = 1, T **do**

        Select action  $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$  according to the current policy and exploration noise

        Execute action  $a_t$  and observe reward  $r_t$  and observe new state  $s_{t+1}$

        Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $R$

        Sample a random minibatch of  $N$  transitions  $(s_i, a_i, r_i, s_{i+1})$  from  $R$

        Set  $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$

        Update critic by minimizing the loss:  $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

        Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

    Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

**end for**

**end for**

---

# Implementation

Suppose we're able to store a *user-item interaction matrix*  $\mathbf{R}$ .



Some cells are empty because there are many items a user do not review ( $\mathbf{R}$  is a **sparse matrix**). But we could fill the empty cells with predictions...

# Implementation

## PROBLEMS:

- Matrix  $R$  is too large
- How do we predict ratings

# Implementation


PROBLEMS:

- Matrix  $R$  is too large
- How do we predict ratings

SOLUTION: **PMF (Probabilistic Matrix Factorization)**


- *Estimating  $R$  by using two low rank matrices  $U$  and  $I$*

We need user and  
item embeddings





# Learn PMF with Neural Networks



$$\text{INTERACTION} = \text{USER\_EMBEDDING} \bullet \text{ITEM\_EMBEDDING}$$

$$\text{PREDICTION} = \text{INTERACTION} + \underbrace{\text{USER\_BIAS} + \text{ITEM\_BIAS}}_{\text{1-dimensional vectors learned by the model with embedding layers}}$$

1-dimensional vectors  
learned by the model with  
embedding layers

# Implementation

Finally we have:

- 1- User embeddings
- 2- Item embeddings
- 3- User x Item rating predictions

WHY DO WE NEED REINFORCEMENT LEARNING?

# Why do we need reinforcement learning

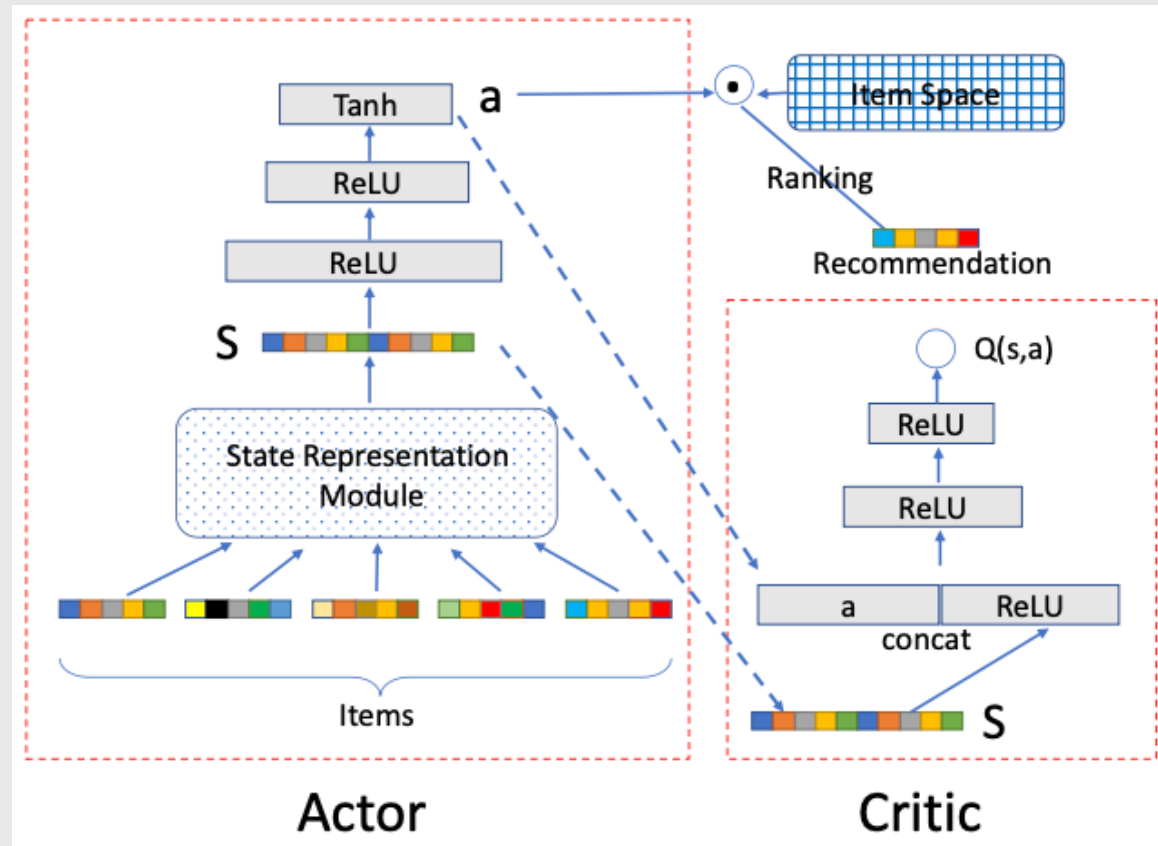
*PMF recommendations are **STATIC**.*

*We want a **DYNAMIC** system that improves recommendations over time according to user preferences.*



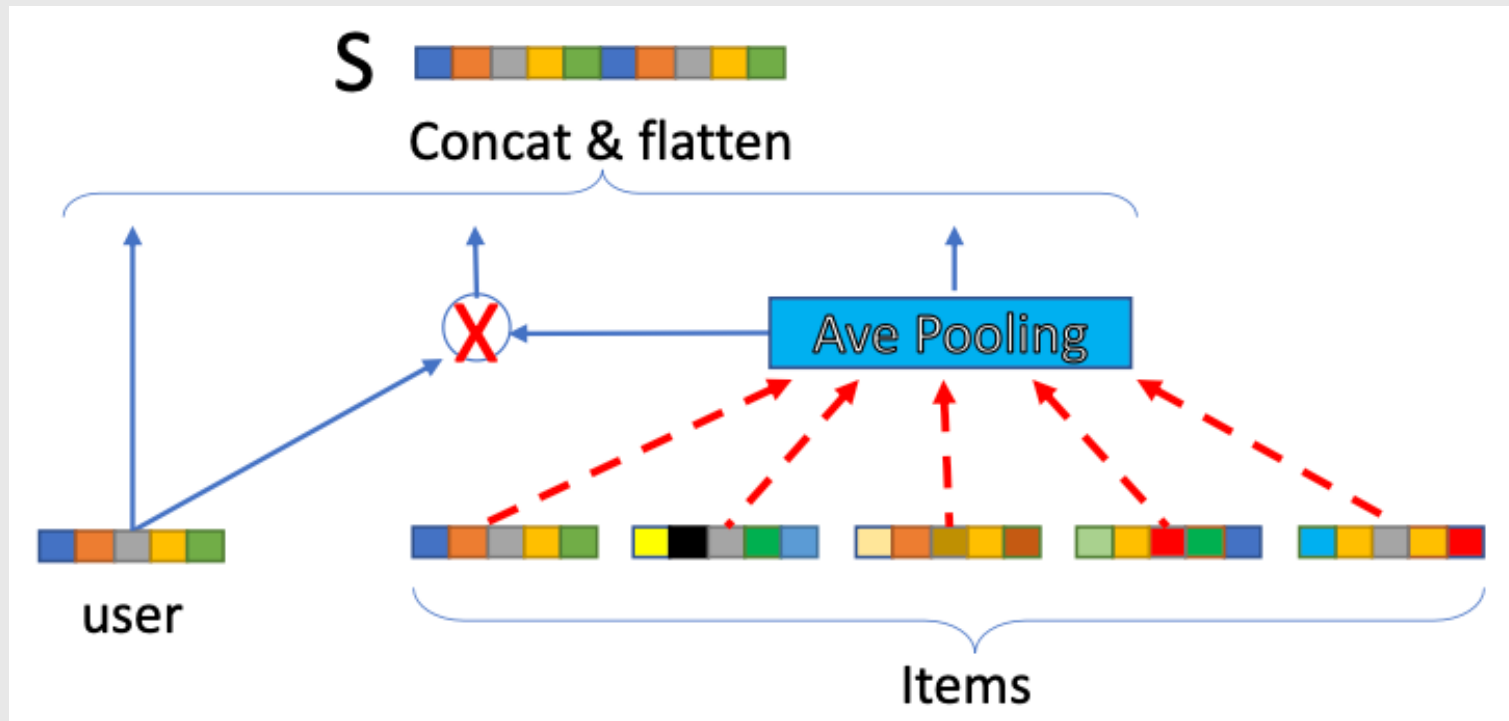
Preferences also change  
over time

# Deep Reinforcement learning based Recommendation (DRR)



# State Representation Module

## DRR-ave




# Two important data structures

- DDR utilizes the user interaction history with the recommender agent as training data.



**We need a HISTORY  
BUFFER H**

- In timestep  $t$ , the training generate a transition  $T$ . Not only we need to store this transition, but we also need to sample a minibatch of  $N$  transitions with prioritized experience sampling technique.



**We need a REPLAY  
BUFFER D**

REMINDER: A transition is a tuple containing state, action, reward and next state given a particular timestep.

# TRAINING ALGORITHM

**input** : Actor learning rate  $\eta_a$ , Critic learning rate  $\eta_c$ , discount factor  $\gamma$ , batch size  $N$ , state window size  $n$  and reward function  $R$

- 1 Randomly initialize the Actor  $\pi_\theta$  and the Critic  $Q_\omega$  with parameters  $\theta$  and  $\omega$
- 2 Initialize the target network  $\pi'$  and  $Q'$  with weights  $\theta' \leftarrow \theta$  and  $\omega' \leftarrow \omega$
- 3 Initialize replay buffer  $D$

$R$  is given by the PMF

After initializing Actor and Critic networks, we simply copied their weights to initialize the target networks

```

4 for  $session = 1, M$  do
5   Observe the initial state  $s_0$  according to the offline
   log
6   for  $t = 1, T$  do
7     Observe current state  $s_t = f(H_t)$ , where
        $H_t = \{i_1, \dots, i_n\}$ 
8     Find action  $a_t = \pi_\theta(s_t)$  according to the current
       policy with  $\varepsilon$ -greedy exploration
9     Recommended item  $i_t$  according to action  $a_t$  by
       Eq. (2)
10    Calculate reward  $r_t = R(s_t, a_t)$  based on the
       feedback of the user
11    Observe new state  $s_{t+1} = f(H_{t+1})$ , where
        $H_{t+1} = \{i_2, \dots, i_n, i_t\}$  if  $r_t$  is positive,
       otherwise,  $H_{t+1} = H_t$ 
12    Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $D$ 
13    Sample a minibatch of  $N$  transitions
        $(s_i, a_i, r_i, s_{i+1})$  in  $D$  with prioritized experience
       replay sampling technique
14    Set  $y_i = r_i + \gamma Q_{\omega'}(s_{i+1}, \pi_{\theta'}(s_{i+1}))$ 
15    Update the Critic network by minimizing the
       loss:  $L = \frac{1}{N} \sum_i (y_i - Q_\omega(s_i, a_i))^2$ 
16    Update the Actor network using the sampled
       policy gradient:
        $\nabla_\theta J(\pi_\theta) \approx$ 
        $\frac{1}{N} \sum_t \nabla_a Q_\omega(s, a)|_{s=s_t, a=\pi_\theta(s_t)} \nabla_\theta \pi_\theta(s)|_{s=s_t}$ 
17    Update the target networks:
        $\theta' \leftarrow \tau\theta + (1 - \tau)\theta'$ 
        $\omega' \leftarrow \tau\omega + (1 - \tau)\omega'$ 
18 return  $\theta$  and  $\omega$ 

```

Each session corresponds  
to a user

$$score_t = i_t a^T$$

Soft update

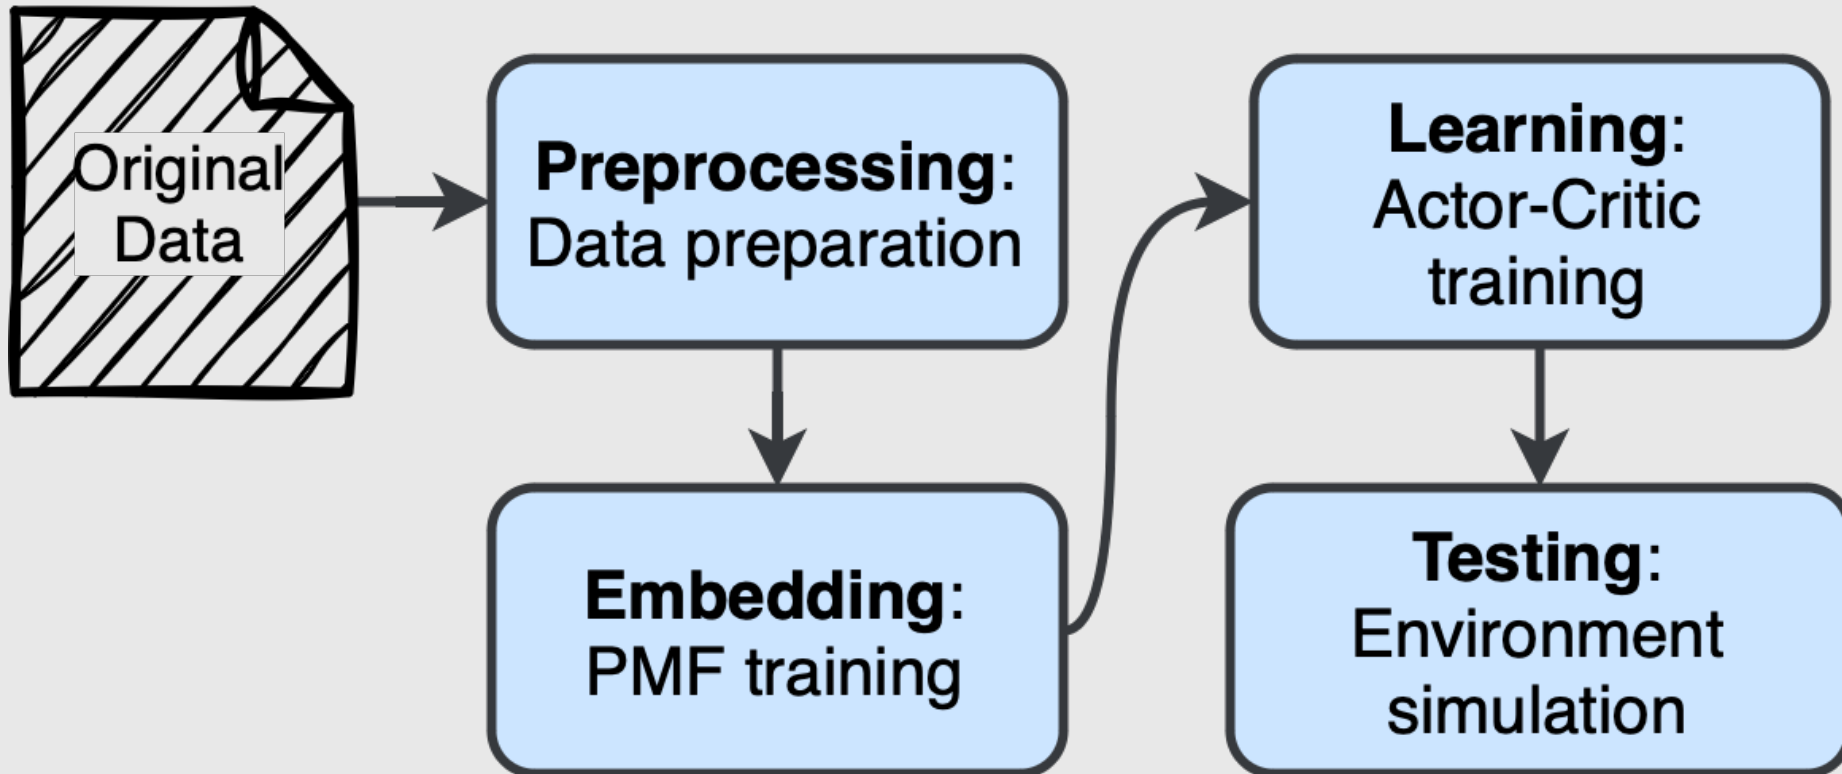


# Experimentation

- **Dataset:** *MovieLens*, ratings by 270000 users for 45000 movies
- **Normalization:** to be used as rewards, ratings  $r \in [1,5]$  have been scaled to  $R^* \in [-1,1]$ , so that positive rewards are given only for  $\geq 3$  rating values:

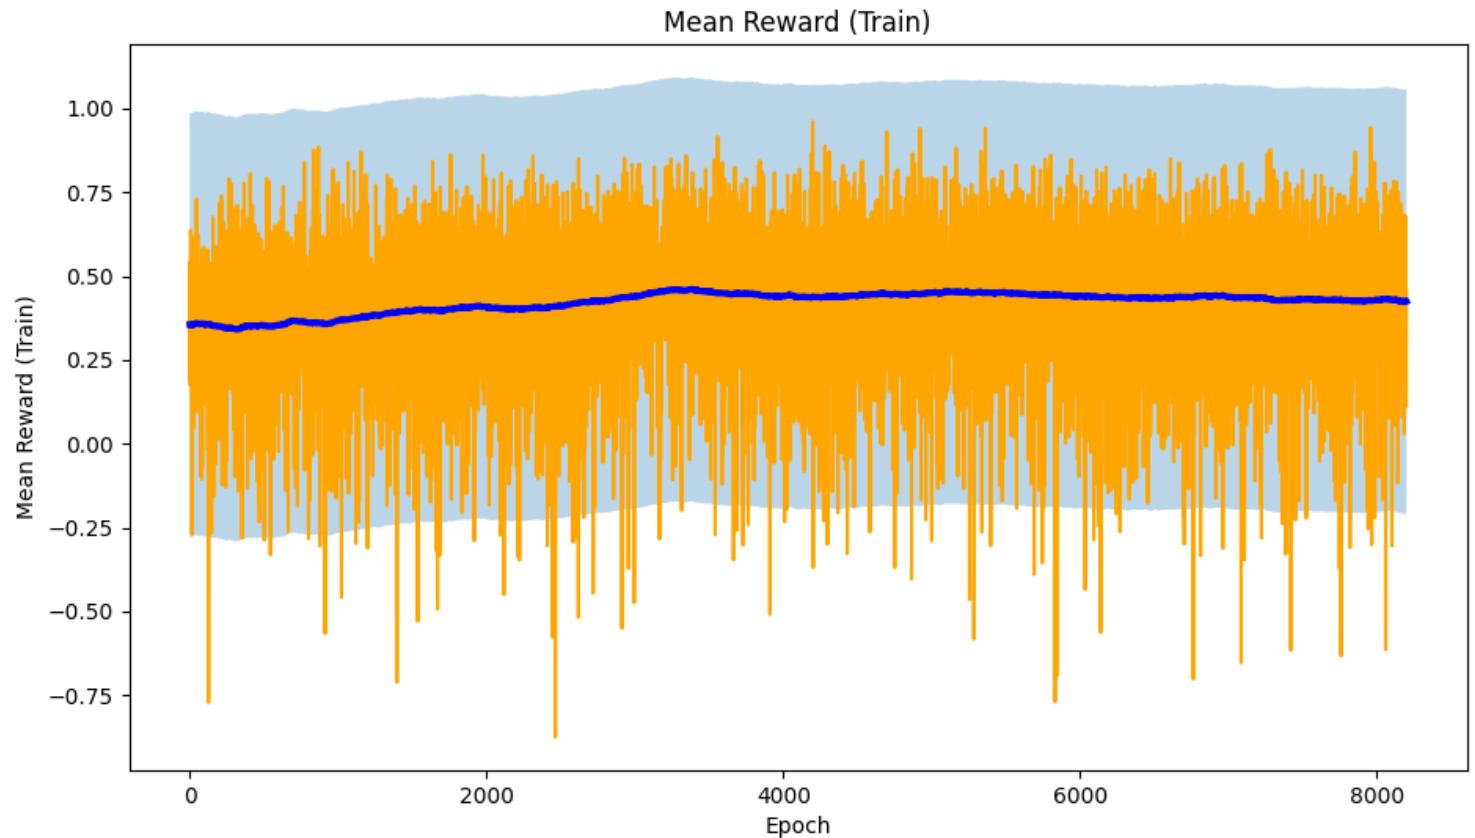

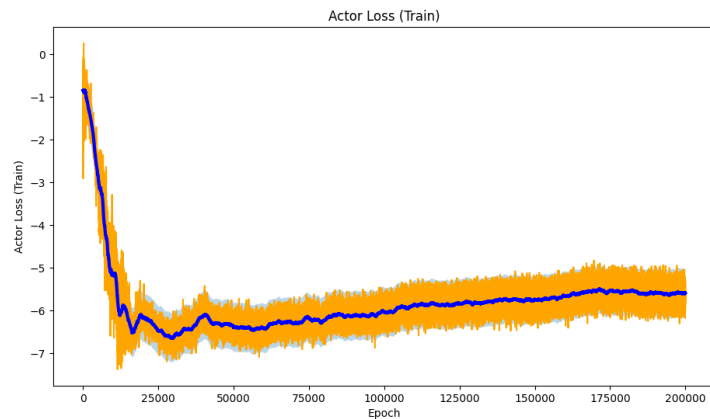
$$R^*(s, a) = \frac{1}{2} (r_i - 3)$$

# Pipeline



# Actor-Critic training

replay buffer = 200000,      history buffer = 5,      lr actor = 0.0005,      lr critic = 0.001  
 $\epsilon_{start} = 1$ ,       $\epsilon_{end} = 0.1$ ,       $\tau = 0.001$ ,       $\beta = 0.4$ ,       $\alpha = 0.3$ ,       $\gamma = 0.9$



# Offline Evaluation

---

**Algorithm 2:** Offline Evaluation Algorithm of DRR Framework

---

**input** : state window size  $n$  and reward function  $R$

- 1 Observe the initial state  $s_0$  and item set  $\mathcal{I}$  according to the offline log
- 2 **for**  $t = 1, T$  **do**
- 3     Observe current state  $s_t = \{i_1, \dots, i_n\}$
- 4     Execute action  $a_t = \pi_\theta(s_t)$  according to the current policy
- 5     Observe the recommended item  $i_t$  according to action  $a_t$  by Eq. (2)
- 6     Get reward  $r_t = R(s_t, a_t)$  from the feedback located in the users' log by Eq. (10)
- 7     Update to a new state  $s_{t+1} = f(H_{t+1})$ , where  $H_{t+1} = \{i_2, \dots, i_n, i_t\}$  if  $r_t$  is positive, otherwise,  $H_{t+1} = H_t$
- 8     remove  $i_t$  from  $\mathcal{I}$

---


For different episode length

$T = [5, 10, 15, 20]$ , in the offline evaluation recommendations are made and judged from the set of items the user has rated rather than all candidate items.

The evaluation is done considering

**Precision @ K**, the amount of positive predictions made divided by the total amount of predictions made ( $K$ ) for a user.


## Trained model



### Trained model precision

@ 5	@ 10	@ 15	@ 20
$0.775 \pm 0.004$	$0.752 \pm 0.003$	$0.733 \pm 0.011$	$0.728 \pm 0.008$

## Randomly initialized model



### Randomly initialized model precision

@ 5	@ 10	@ 15	@ 20
$0.672 \pm 0.004$	$0.665 \pm 0.010$	$0.655 \pm 0.004$	$0.649 \pm 0.004$

### Paper precision results

@ 5	@ 10
0.7693	0.6594

# Online Evaluation

---

**Algorithm 1:** Training Algorithm of DRR Framework

---


**input** : Actor learning rate  $\eta_a$ , Critic learning rate  $\eta_c$ , discount factor  $\gamma$ , batch size  $N$ , state window size  $n$  and reward function  $R$

- 1 Randomly initialize the Actor  $\pi_\theta$  and the Critic  $Q_\omega$  with parameters  $\theta$  and  $\omega$
- 2 Initialize the target network  $\pi'$  and  $Q'$  with weights  $\theta' \leftarrow \theta$  and  $\omega' \leftarrow \omega$
- 3 Initialize replay buffer  $D$
- 4 **for**  $session = 1, M$  **do**
- 5   Observe the initial state  $s_0$  according to the offline log
- 6   **for**  $t = 1, T$  **do**
- 7     Observe current state  $s_t = f(H_t)$ , where  $H_t = \{i_1, \dots, i_n\}$
- 8     Find action  $a_t = \pi_\theta(s_t)$  according to the current policy with  $\varepsilon$ -greedy exploration
- 9     Recommended item  $i_t$  according to action  $a_t$  by Eq. (2)
- 10    Calculate reward  $r_t = R(s_t, a_t)$  based on the feedback of the user
- 11    Observe new state  $s_{t+1} = f(H_{t+1})$ , where  $H_{t+1} = \{i_2, \dots, i_n, i_t\}$  if  $r_t$  is positive, otherwise,  $H_{t+1} = H_t$
- 12    Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $D$
- 13    Sample a minibatch of  $N$  transitions  $(s_i, a_i, r_i, s_{i+1})$  in  $D$  with prioritized experience replay sampling technique
- 14    Set  $y_i = r_i + \gamma Q_{\omega'}(s_{i+1}, \pi_{\theta'}(s_{i+1}))$
- 15    Update the Critic network by minimizing the loss:  $L = \frac{1}{N} \sum_i (y_i - Q_\omega(s_i, a_i))^2$
- 16    Update the Actor network using the sampled policy gradient:  
$$\nabla_\theta J(\pi_\theta) \approx \frac{1}{N} \sum_t \nabla_a Q_\omega(s, a)|_{s=s_t, a=\pi_\theta(s_t)} \nabla_\theta \pi_\theta(s)|_{s=s_t}$$
- 17    Update the target networks:  
$$\theta' \leftarrow \tau \theta + (1 - \tau) \theta'$$
$$\omega' \leftarrow \tau \omega + (1 - \tau) \omega'$$
- 18 **return**  $\theta$  and  $\omega$

---

For different episode length  $T = [5, 10, 15, 20]$ , the online evaluation follows the training algorithm with the only difference being the networks and their parameters are reset to those of the learned networks before each episode. This is close to a real scenario where users start with the base model and continue training or refining the model with their own interactions. The evaluation is done considering **Avg.Reward @ K**, the average reward per user receiving K recommendations.


## Trained model



Trained model average reward

@ 5	@ 10	@ 15	@ 20
$0.662 \pm 0.004$	$0.629 \pm 0.004$	$0.594 \pm 0.012$	$0.581 \pm 0.008$





## Randomly initialized model



Randomly initialized average reward



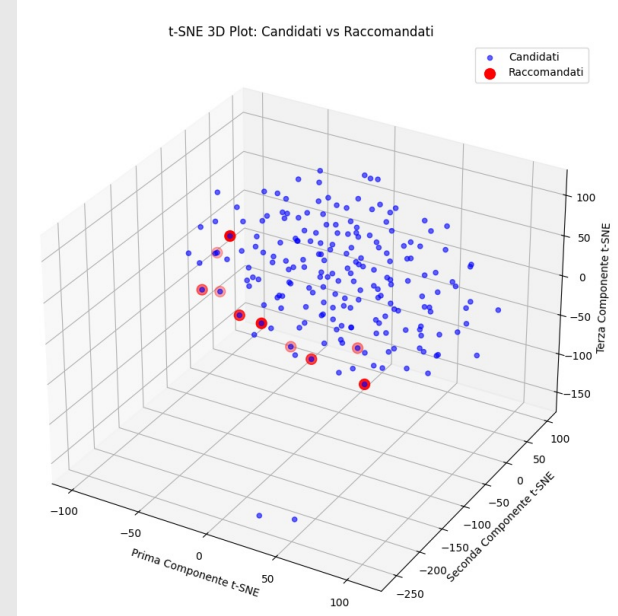

@ 5	@ 10	@ 15	@ 20
$0.323 \pm 0.005$	$0.381 \pm 0.009$	$0.392 \pm 0.011$	$0.397 \pm 0.009$

# Recommended items visualization





# Recommended items visualization



# CONCLUSIONS